

ECE521 Inference Algorithms and Machine Learning

Report of Assignment 3

Group Members:

Xi Chen (Student ID: 1002266013)

Yang Wu (Student ID: 1002495132)

Task 1. K-means

1.1 Euclidean distance function

1.

$$\begin{aligned}
 D(X, Y) &= \begin{bmatrix} D(X_1, Y_1) & D(X_1, Y_2) & \cdots & D(X_1, Y_K) \\ D(X_2, Y_1) & D(X_2, Y_2) & \cdots & D(X_2, Y_K) \\ \vdots & \vdots & \ddots & \vdots \\ D(X_B, Y_1) & D(X_B, Y_2) & \cdots & D(X_B, Y_K) \end{bmatrix}_{B \times K} \\
 &= \begin{bmatrix} (X_1 - Y_1)(X_1 - Y_1)^T & (X_1 - Y_2)(X_1 - Y_2)^T & \cdots & (X_1 - Y_K)(X_1 - Y_K)^T \\ (X_2 - Y_1)(X_2 - Y_1)^T & (X_2 - Y_2)(X_2 - Y_2)^T & \cdots & (X_2 - Y_K)(X_2 - Y_K)^T \\ \vdots & \vdots & \ddots & \vdots \\ (X_B - Y_1)(X_B - Y_1)^T & (X_B - Y_2)(X_B - Y_2)^T & \cdots & (X_B - Y_K)(X_B - Y_K)^T \end{bmatrix}_{B \times K} \\
 &= \begin{bmatrix} X_1 X_1^T & X_1 X_1^T & \cdots & X_1 X_1^T \\ X_2 X_2^T & X_2 X_2^T & \cdots & X_2 X_2^T \\ \vdots & \vdots & \ddots & \vdots \\ X_B X_B^T & X_B X_B^T & \cdots & X_B X_B^T \end{bmatrix}_{B \times K} + \begin{bmatrix} Y_1 Y_1^T & Y_2 Y_2^T & \cdots & Y_K Y_K^T \\ Y_1 Y_1^T & Y_2 Y_2^T & \cdots & Y_K Y_K^T \\ \vdots & \vdots & \ddots & \vdots \\ Y_1 Y_1^T & Y_2 Y_2^T & \cdots & Y_K Y_K^T \end{bmatrix}_{B \times K} - 2 \begin{bmatrix} X_1 Y_1^T & X_1 Y_2^T & \cdots & X_1 Y_K^T \\ X_2 Y_1^T & X_2 Y_2^T & \cdots & X_2 Y_K^T \\ \vdots & \vdots & \ddots & \vdots \\ X_B Y_1^T & X_B Y_2^T & \cdots & X_B Y_K^T \end{bmatrix}_{B \times K}
 \end{aligned}$$

Where

- The first term can be derived from:

$$XX^T = \begin{bmatrix} X_1 X_1^T & X_1 X_2^T & \cdots & X_1 X_B^T \\ X_2 X_1^T & X_2 X_2^T & \cdots & X_2 X_B^T \\ \vdots & \vdots & \ddots & \vdots \\ X_B X_1^T & X_B X_2^T & \cdots & X_B X_B^T \end{bmatrix}_{B \times K} ; A = \text{diag}(XX^T) = \begin{bmatrix} X_1 X_1^T & 0 & \cdots & 0 \\ 0 & X_2 X_2^T & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & X_B X_B^T \end{bmatrix}_{B \times K}$$

$$\text{Hence, } A \cdot \begin{bmatrix} 1 & 1 & \cdots & 1 \\ 1 & 1 & \cdots & 1 \\ \vdots & \vdots & \ddots & \vdots \\ 1 & 1 & \cdots & 1 \end{bmatrix}_{B \times K} = \begin{bmatrix} X_1 X_1^T & X_1 X_1^T & \cdots & X_1 X_1^T \\ X_2 X_2^T & X_2 X_2^T & \cdots & X_2 X_2^T \\ \vdots & \vdots & \ddots & \vdots \\ X_B X_B^T & X_B X_B^T & \cdots & X_B X_B^T \end{bmatrix}_{B \times K}$$

$$\text{Therefore, } D(X, Y)_1 = \text{diag}(XX^T) \cdot [1]_{B \times K} = \text{diag}(XX^T) \cdot I_{B \times K}$$

- The second term can be derived in the similar process:

$$D(X, Y)_2 = (\text{diag}(YY^T) \cdot [1]_{K \times B})^T = (\text{diag}(YY^T) \cdot I_{K \times B})^T = I_{K \times B}^T \cdot \text{diag}(YY^T)^T$$

- The third term is simply as:

$$XY^T = \begin{bmatrix} X_1 \\ X_2 \\ \vdots \\ X_B \end{bmatrix}_{B \times D} \cdot [Y_1^T \quad Y_2^T \quad \cdots \quad Y_K^T]_{D \times K} = \begin{bmatrix} X_1 Y_1^T & X_1 Y_2^T & \cdots & X_1 Y_K^T \\ X_2 Y_1^T & X_2 Y_2^T & \cdots & X_2 Y_K^T \\ \vdots & \vdots & \ddots & \vdots \\ X_B Y_1^T & X_B Y_2^T & \cdots & X_B Y_K^T \end{bmatrix}_{B \times K}$$

$$D(X, Y)_3 = 2 * XY^T$$

Therefore, $D(X,Y) = D(X,Y)_1 + D(X,Y)_2 - D(X,Y)_3$

$$= \text{diag}(XX^T) \cdot I_{B \times K} + I_{K \times B}^T \cdot \text{diag}(YY^T)^T - 2 * XY^T$$

2. The vectorized function has been written in the file "**Euclid_Distance.py**", which calculates Euclidean distance between possible pairs of vector points. The first two terms we derived above can be implemented by using Tensorflow broadcasting. The code is also listed as below:

```
class Euclid_Distance:
    def __init__(self, X, Y, D):
        self.X = X
        self.Y = Y
        self.D = D

    def cal_Euclid_dis(self):
        x2 = self.cal_square(self.X)
        y2 = self.cal_square(self.Y)
        xy = self.cal_XY(self.X, self.Y)

        Euclid_dist = x2 + tf.transpose(y2) - 2*xy
        return Euclid_dist

    def cal_square(self, X):
        square = tf.square(X)
        result = tf.matmul(square, tf.ones(shape=[self.D, 1]))
        return result

    def cal_XY(self, X, Y):
        result = tf.matmul(X, Y, False, True)
        return result
```

Fig1. Implements the pair-wise squared Euclidean distance function for two input matrices

1.2 Learning K-means

1. The loss function $L(\mu)$ is non-convex. If the function was convex, it would always converge to the global minimum; hence, the final value of $L(\mu)$ would be fixed under any circumstances. In our experiment, however, the loss function has different local minima and the final result depends on how the K cluster centers are initialized. The final result is only guaranteed to be one of these local minima, but it can be difficult to achieve the global minimum. Therefore, the loss function $L(\mu)$ is a non-convex function.
2. The functions are implemented in "**k_mean.py**". We set the learning rate as 0.001 and get the cluster centers are:

$$\begin{aligned}\mu^1 & [1.24792147, 0.25577992] \\ \mu^2 & [0.15492232, -1.51739335] \\ \mu^3 & [-1.05739701, -3.23398256]\end{aligned}$$

Fig2. shows the loss vs. number updates.





Fig2. The loss Vs. the number of updates for K-mean (K = 3)


We can see from the Fig1. that the loss function will converge to a local minimum after a few hundreds updates.

3. We run the Algorithm with K = 1, 2, 3, 4 and 5. The 2D scatter of the data points and the percentage of the data point belonging to each cluster are listed below:




K = 2

The percentages of the data points belonging to the two clusters are **50.46%** and **49.54%**.




K = 3

The percentage of the data points belonging to the three clusters are **38.13%**, **23.81%** and **38.06%**.



K = 4

The percentage of the data points belonging to the four clusters are **37.13%**, **13.53%**, **37.30%** and **12.04%**.






Fig3. 2D scatter plot of data points colored by their cluster assignments

We can see from the above plot that when all data points are divided by two clusters, either side has almost the equal percentage (~50%). It also emphasized the symmetric characteristic of the plot. When K continues increasing, the centroids of two clusters with larger data density remain almost unchanged. While the central part of the data points is divided into the extra clusters according to the distance from data to the centroids. Thus, we think the central part of data points can be labeled as one cluster, and the rest part should be the other two clusters.

Thus, we think if we want to divide the data into clusters with equal number points, $k = 2$ gives the best result; otherwise, $k = 3$ should be the best value for clustering.

4. Our result shows that the loss function for validation data is minimum when k equals to 5. Thus, $K = 5$ is the best value if we want to minimize the loss function for validation data. Actually, the more clusters there are, the less the loss will be. It will be the best scenario if there exist 10000 clusters, so each cluster is corresponding to a data point. In this case, every cluster has the zero loss, and the total one will be zero as well. If K continues growing, the additional cluster will be useless.

Here are the losses for the validation data when k ranges from 1 to 5:

K	Loss for validation data
1	12969.2
2	3065.8
3	1693.22
4	1100.41
5	934.067

Task 2. Mixtures of Gaussians

2.1 The Gaussian cluster model

1.The expression for the latent variable posterior distribution $P(z|x)$ is:

$$P(z = k|x) = \frac{P(x|z = k)P(z = k)}{\sum_{z'=1}^K P(x|z')P(z')} = \frac{\mathcal{N}(x; \mu^k, \sigma^{k^2})\pi^k}{\sum_{j=1}^K \mathcal{N}(x; \mu^j, \sigma^{j^2})\pi^j} = \frac{\frac{\pi^k}{\sqrt{(2\pi\sigma^{k^2})^D}} e^{-\frac{(x-\mu^k)^2}{2\sigma^{k^2}}}}{\sum_{j=1}^K \left(\frac{\pi^j}{\sqrt{(2\pi\sigma^{j^2})^D}} e^{-\frac{(x-\mu^j)^2}{2\sigma^{j^2}}} \right)}$$

where σ is the standard deviation, μ is the mean and D is the dimension of dataset.

2.The log probability density function for cluster k is:

$$\log \mathcal{N}(x; \mu^k, \sigma^{k^2}) = \log \frac{1}{\sqrt{(2\pi\sigma^{k^2})^D}} e^{-(x-\mu^k)^2/2\sigma^{k^2}} = -\frac{D}{2} \log(2\pi\sigma^{k^2}) - \frac{(x-\mu^k)^2}{2\sigma^{k^2}}$$

The function is implemented in the “**Log_Probability.py**” file. To get the log probability for all B data points and K clusters, we use the Tensorflow broadcasting. The code is also listed as below:

```
class Log_Probability:
    def __init__(self, X, Y, sigma, D):
        self.X = X
        self.Y = Y
        self.sigma = sigma
        self.D = D

    def cal_Euclid_dis(self):
        x2 = self.cal_square(self.X)
        y2 = self.cal_square(self.Y)
        xy = self.cal_XY(self.X, self.Y)

        Euclid_dist = x2 + tf.transpose(y2) - 2*xy
        return Euclid_dist

    def cal_square(self, X):
        square = tf.square(X)
        result = tf.matmul(square, tf.ones(shape=[self.D, 1], dtype=tf.float32))
        return result

    def cal_XY(self, X, Y):
        result = tf.matmul(X, Y, False, True)
        return result

    def cal_Term1(self, sigma):
        return -(self.D/2) * tf.log(2 * pi * tf.square(self.sigma))

    def cal_Term2(self, ed, sigma):
        return tf.div(-ed, 2 * tf.square(self.sigma))

    def cal_log_probability(self):
        ed = self.cal_Euclid_dis()
        log_prob = self.cal_Term1(self.sigma) + self.cal_Term2(ed, self.sigma)
        return log_prob
```

Fig4. Implements of log probability density functions for cluster K

3.The log probability of the cluster variable z given the data vector \mathbf{x} , $P(z|\mathbf{x})$ is:

$$\log P(z = k|\mathbf{x}) = P(\mathbf{x}, z = k) / \sum_{z'=1}^K P(\mathbf{x}, z')$$

$$\begin{aligned}
 &= \log \frac{\pi^k \cdot \frac{1}{\sqrt{(2\pi\sigma^{k^2})^D}} e^{-(\mathbf{x}-\boldsymbol{\mu}^k)^2/2\sigma^{k^2}}}{\sum_{k'=1}^K \left(\pi^{k'} \cdot \frac{1}{\sqrt{(2\pi\sigma^{k'^2})^D}} e^{-(\mathbf{x}-\boldsymbol{\mu}^{k'})^2/2\sigma^{k'^2}} \right)} \\
 &= \log \frac{e^{\left(-\frac{(\mathbf{x}-\boldsymbol{\mu}^k)^2}{2\sigma^{k^2}} + \log \frac{\pi^k}{\sqrt{(2\pi\sigma^{k^2})^D}} \right)}}{\sum_{k'=1}^K e^{\left(-\frac{(\mathbf{x}-\boldsymbol{\mu}^{k'})^2}{2\sigma^{k'^2}} + \log \frac{\pi^{k'}}{\sqrt{(2\pi\sigma^{k'^2})^D}} \right)}} \\
 &= -\frac{(\mathbf{x}-\boldsymbol{\mu}^k)^2}{2\sigma^{k^2}} + \log \frac{\pi^k}{\sqrt{(2\pi\sigma^{k^2})^D}} - \log \sum_{k'=1}^K e^{\left(-\frac{(\mathbf{x}-\boldsymbol{\mu}^{k'})^2}{2\sigma^{k'^2}} + \log \frac{\pi^{k'}}{\sqrt{(2\pi\sigma^{k'^2})^D}} \right)}
 \end{aligned}$$

where $-\frac{(\mathbf{x}-\boldsymbol{\mu}^k)^2}{2\sigma^{k^2}} + \log \frac{\pi^k}{\sqrt{(2\pi\sigma^{k^2})^D}}$ is the input tensor of `reduce_logsumexp()` function. The

function is implemented in the “**Log_Posterior.py**” file. The code is also listed below:

```

class Log_Posterior:
    def __init__(self, X, Y, sigma, pi_k, D):
        self.X = X
        self.Y = Y
        self.sigma = sigma
        self.pi_k = pi_k
        self.D = D

    def cal_Euclid_dis(self):
        x2 = self.cal_square(self.X)
        y2 = self.cal_square(self.Y)
        xy = self.cal_XY(self.X, self.Y)

        Euclid_dist = x2 + tf.transpose(y2) - 2*xy
        return Euclid_dist

    def cal_square(self, X):
        square = tf.square(X)
        result = tf.matmul(square, tf.ones(shape=[self.D, 1], dtype=tf.float32))
        return result

    def cal_XY(self, X, Y):
        result = tf.matmul(X, Y, False, True)
        return result

    def cal_term1(self, pi_k, sigma):
        return tf.log(tf.div(pi_k, tf.sqrt(tf.pow(2 * pi * tf.square(sigma), self.D))))

    def cal_term2(self, sigma):
        ed = self.cal_Euclid_dis()
        return tf.mul(-0.5, tf.div(ed, tf.square(sigma)))

    def cal_term3(self):
        my_tensor = self.cal_term2(self.sigma) + self.cal_term1(self.pi_k, self.sigma)
        log_sum = reduce_logsumexp(my_tensor, 1, True)
        return log_sum

    def cal_log_posterior(self):
        res = self.cal_term1(self.pi_k, self.sigma) + self.cal_term2(self.sigma) - self.cal_term3()
        return res

```

Fig5. Implementations of log probability of the cluster variable z given the data vector x

We note that the input tensor may be too small when the standard deviation is very small, such that the exponentiation is 0. When we use the `tf.reduce_sum`, the return of logarithm could be `-inf`. Thus, we should use **reduce_logsumexp** here, to ensure that we shift the center of the exponentiated variables, such that the largest value we want to exponentiate is zero. By doing this change, even the rest exponentiations will lead to underflow, we could still get a reasonable value. The whole process is listed as following:

$$\begin{aligned}
 \log \sum_{k=1}^K e^{x_k} &= \log(e^{x_1} + e^{x_2} + \dots + e^{x_k}) \\
 &= \log(e^{x_1-x_m} + e^{x_2-x_m} + \dots + e^0 + \dots + e^{x_k-x_m})e^{x_m} \\
 &= \log(e^{x_1-x_m} + e^{x_2-x_m} + \dots + e^0 + \dots + e^{x_k-x_m}) + x_m
 \end{aligned}$$

where the x_m is the maximum value in the input tensor. Even when x_k (for $k = 1$ to K) is very small, we still can get a reasonable result in this way.

2.2 Learning the MoG

1.

$$\begin{aligned}
 \nabla \log P(x) &= \nabla \log \sum_{k=1}^K P(z=k)P(x|z=k) \\
 &= \frac{\nabla \sum_{k=1}^K P(z=k)P(x|z=k)}{\sum_{k'=1}^K P(z=k')P(x|z=k')} \\
 &= \frac{\sum_{k=1}^K \nabla P(z=k)P(x|z=k)}{\sum_{k'=1}^K P(z=k')P(x|z=k')} \\
 &= \frac{\sum_{k=1}^K P(z=k)P(x|z=k) \nabla \log P(z=k)P(x|z=k)}{\sum_{k'=1}^K P(z=k')P(x|z=k')} \\
 &= \sum_{k=1}^K \left(\frac{P(z=k)P(x|z=k)}{\sum_{k'=1}^K P(z=k')P(x|z=k')} \right) \nabla \log P(z=k)P(x|z=k) \\
 &= \sum_k P(z=k|x) \nabla \log P(x, z=k)
 \end{aligned}$$

2. We set the learning rate of AdamOptimizer as 0.01 and use the gradient descent to minimize the negative log likelihood. The function is implemented in “**MoG.py**”.

When K = 3, we get these model parameters:

μ^k	$\mu^1[0.10617115, -1.52815366]$ $\mu^2[-1.10174823, -3.30640173]$ $\mu^3[1.298244, 0.30914176]$
π^k	[0.33480254, 0.33164248, 0.33355504]
σ^{k^2}	[0.98666757, 0.03905554, 0.03885973]

The loss Vs. the number of updates are shown in Fig6.




Fig6. The loss Vs. the number of updates for MoG (K = 3)


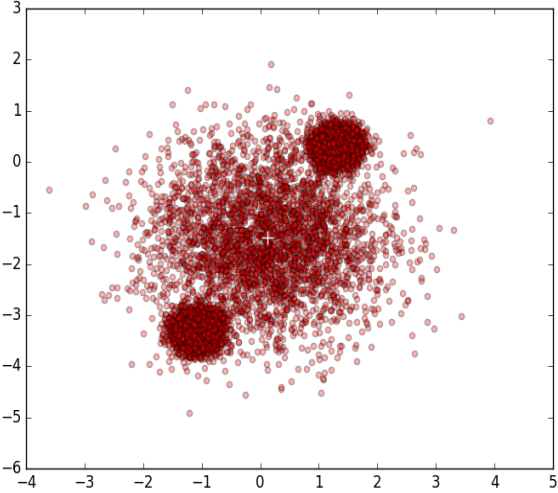
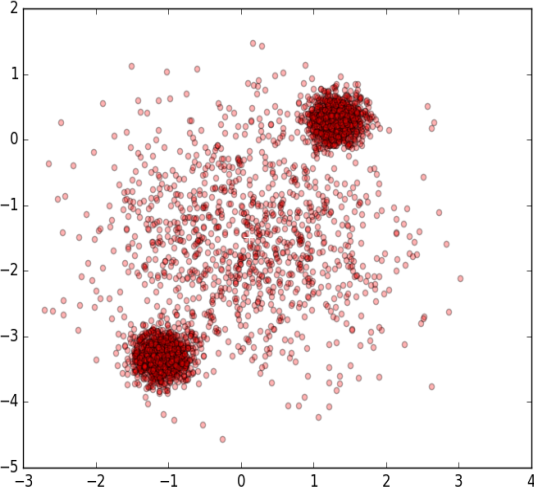
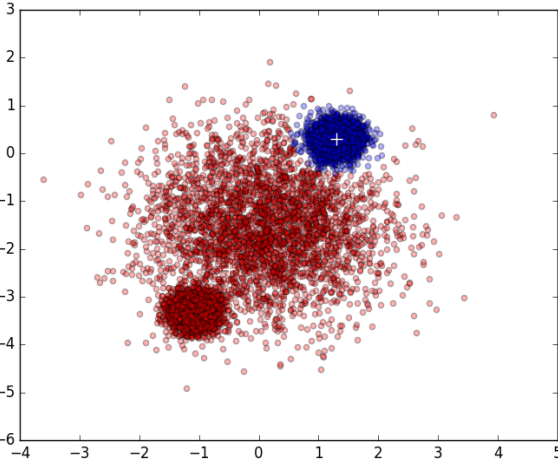
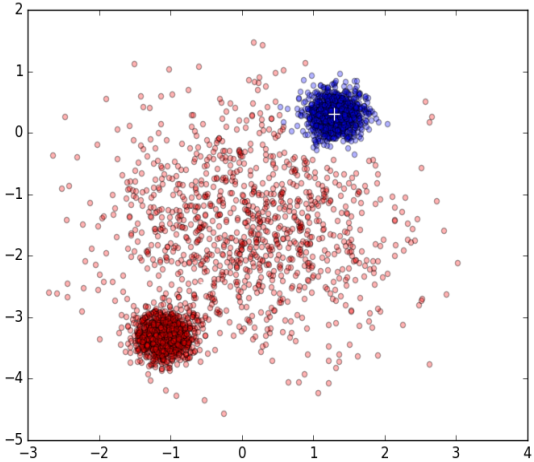


Fig7. 2D scatter of data points (K = 3)

We could see from Fig3 and the model parameters that the loss function converges to 17000 after about 500 times updates. The data points are divided into three clusters. One has a larger standard deviation and the other two clusters have a much smaller standard deviation.

3. When 1/3 of the data is hold for validation and set K from 1 to 5, the MoG mode give us the following result.

	Complete data	Validation data
<div><div>K = 1</div><div>Loss for validation data is 11680.0.</div></div>		
<div><div>K = 2</div><div>Loss for validation data is 8051.42.</div></div>		

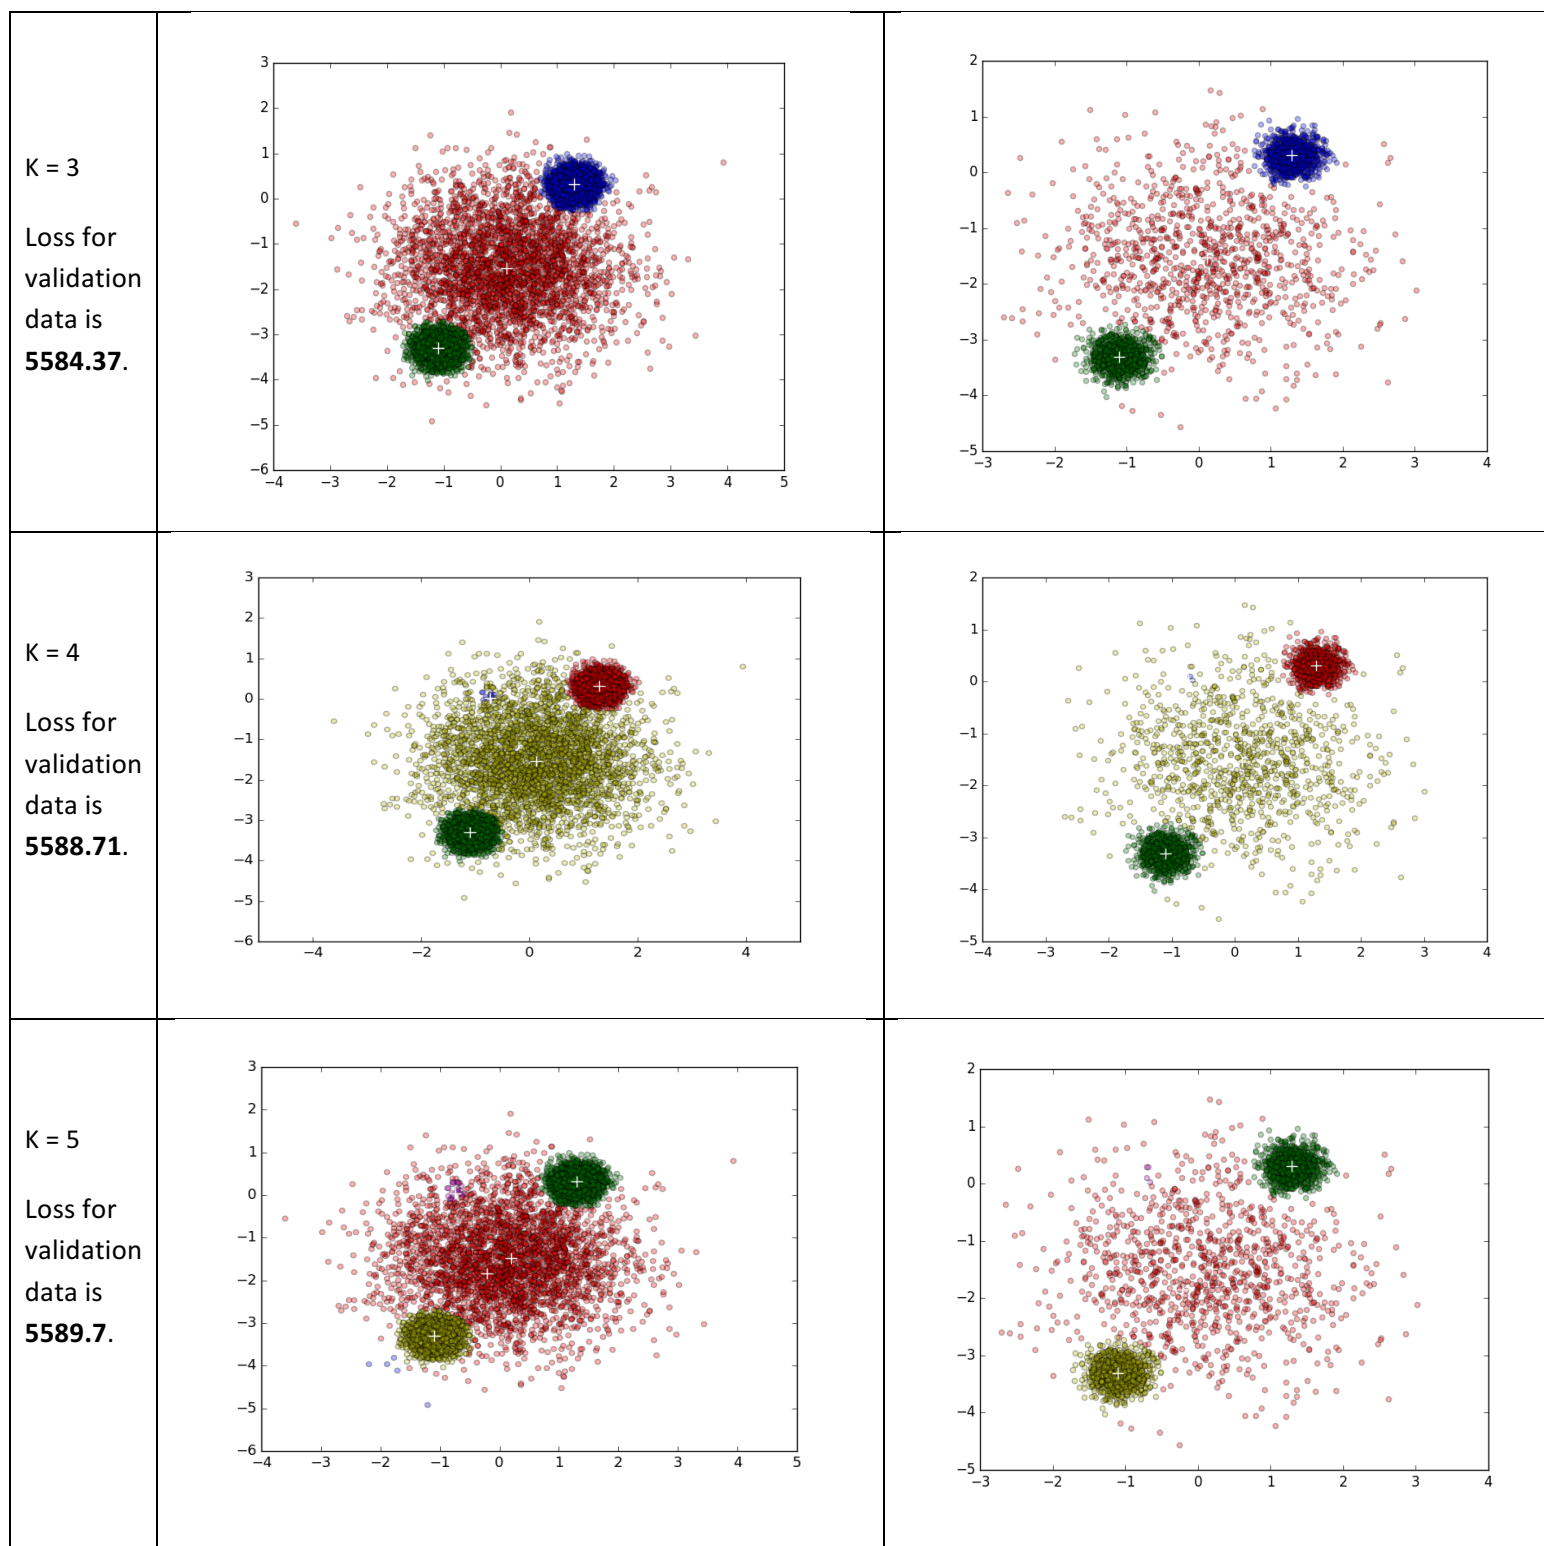


Fig8. 2D scatter plot of data points by using MoG for different K values
(Complete data and validation data)

From the plots and loss function results for validation data, we could conclude that **3** is the best value for clusters to divide the data points. **When K = 3, we receive the minimum loss** and when K continues increasing, the loss will increase a little.

In the scatter plots of data points colored by cluster assignments, we could observe that when K = 3, the data points are evenly divided into three clusters. However, when k increases to 4 and 5, there will be a small number of data points (~10-20) to be assigned to extra clusters. This extra clusters have much smaller π^k , which have little effect on the clustering. Thus, the likelihood is extremely small for data points to be assigned to these extra clusters. By training the MoG model for lots of times, we find that the percentage of data points to be assigned to the fourth and fifth cluster is really small (e.g. 0.001), sometimes even to be zero. Thus, **we think K = 3 is the best value for the data points by using the MoG method.**

4. When we run both K-means and MoG on the data100D.npy, we got the following result:

Proportion for different clusters in k-means:

(In all these cases, we try many times to get the global minimum, otherwise there could be only 3 clusters when k = 4 and 4 clusters when k = 5.)

K	Percentage
1	[1.]
2	[0.20, 0.80]
3	[0.20, 0.40, 0.40]
4	[0.30, 0.20, 0.30, 0.20]
5	[0.10, 0.20, 0.20, 0.30, 0.20]

Parameters and proportions for different clusters in MoG:

(We also try many times for each value of K, but each time we can only get 1 cluster by MoG.)

K = 1:

π^k	[1.]
σ^2	[1.00000036]
Cluster percentage	[1.]

K = 2:

π^k	[0.5 0.5]
σ^2	[0.78500676 0.80555862]
Cluster percentage	[1. 0.]

K = 3:

π^k	[0.5, 0.2, 0.3]
σ^2	[0.68058121, 0.48646179, 0.57936233]
Cluster percentage	[1. 0. 0.]

K = 4:

π^k	[4.99916792e-01, 1.89150320e-04, 1.99961022e-01, 2.99932897e-01]
σ^2	[0.62175423, 0.42206311, 0.12680988, 0.78740317]
Cluster percentage	[1. 0. 0. 0.]

K = 5:

π^k	[0.10002382, 0.2000609, 0.20001048, 0.19994211, 0.29996258]
σ^2	[0.01266976, 0.12677838, 0.4864502, 0.09153681, 0.24339898]
Cluster percentage	[1. 0. 0. 0. 0]

We run the two methods many times, and find that K-means can divide the 100-dimensions data into K clusters but with different percentages of data points for different variable initializations. The different percentage results from the random initialization of centroids.

But for the MoG method, there is only 1 cluster as we try many times, each time we could get similar model parameters. Thus, we think each data point has been labeled to the cluster with the maximum likelihood.

Therefore, we think there is only **one** cluster within the dataset.